

Building a Kubernetes Operator with Kubebuilder

1. Pre-requisites

Before we start, make sure you have the following installed:

- Go (v1.13 or higher)
- Docker (for building and pushing the operator image)
- **kubect1** (to interact with your cluster)
- A Kubernetes cluster (like Minikube or KIND)

2. Install Kubebuilder

Install the latest release of **kubebuilder** from the [official release page](#).

3. Initialize Your Operator Project

Let's start by creating a new directory for our operator and navigating into it:

```
mkdir -p /mnt/c/project/nimble-opti-adapter  
cd /mnt/c/project/nimble-opti-adapter
```

Now, initialize the project:

```
kubebuilder init --domain uri-tech.github.io --repo github.com/uri-  
tech/nimble-opti-adapter
```

4. Define Your Custom Resource (CR)

CRs are the heart of your operator. Let's define one:

```
kubebuilder create api --group=adapter --version=v1 --kind=NimbleOpti --  
resource=true --controller=true
```

This will scaffold the **NimbleOpti** CRD and its associated controller.

5. Implement the Operator Logic

Navigate to the generated controller file, likely at:

controllers/adapter_nimbleopti_controller.go

Here, implement the reconciliation logic inside the **Reconcile** method. This is where the magic happens!

6. Watch for Ingress and NimbleOpti Resources

To ensure our operator reacts to changes in Ingress and NimbleOpti resources, add watches for these types in the controller's setup.

Modify the `SetupWithManager` method in `controllers/adapter_nimbleopti_controller.go` to include:

```
// SetupWithManager sets up the controller with the Manager. The Manager
// will set fields on the Controller
// and Start it when the Manager is Started.
func (r *NimbleOptiReconciler) SetupWithManager(mgr ctrl.Manager) error {
    // debug
    klog.Infof("debug - SetupWithManager")

    // NewControllerManagedBy returns a builder for a controller
    // managed by mgr that will be started with mgr.Start.
    b := ctrl.NewControllerManagedBy(mgr)

    // For the primary resource type that this controller watches
    b = b.For(&adaperv1.NimbleOpti{})

    // Owns specifies objects that are owned by the primary resource
    // The argument here must be a runtime object that will have its
    // Group, Version, and Kind filled in.
    b = b.Owns(&networkingv1.Ingress{})

    // WithEventFilter specifies a Predicate that will be used to filter
    // events before they are sent to event handlers.
    // The GenerationChangedPredicate filters out objects that have not
    // changed their .metadata.generation field.
    // b = b.WithEventFilter(predicate.GenerationChangedPredicate{})

    // Watch Ingress objects – do not need it – do it from the
    IngressWatcher
    // if err := b.Watches(&networkingv1.Ingress{},
    &handler.EnqueueRequestForObject{}); err != nil {
    //     klog.Error(err, "unable to watch Ingress")
    // }

    // Watch for changes to Pods
    // if err := b.Watches(&source.Kind{Type: &corev1.Pod{}},
    &handler.EnqueueRequestForObject{}); err != nil {
    //     klog.Error(err, "unable to watch Pods")
    // }

    // Call Complete to create the NimbleOptiReconciler. This step comes
    at the end
    // as it finalizes the controller's configuration.
    return b.Complete(r)
}
```

7. Update the operator core files:

./cmd/main.go:

```
package main

import (
    "flag"
    "os"

    // Import all Kubernetes client auth plugins (e.g. Azure, GCP, OIDC,
    etc.)
    // to ensure that exec-entrypoint and run can make use of them.
    "k8s.io/client-go/kubernetes"
    _ "k8s.io/client-go/plugin/pkg/client/auth"
    "k8s.io/klog/v2"

    "k8s.io/apimachinery/pkg/runtime"
    utilruntime "k8s.io/apimachinery/pkg/util/runtime"
    clientgoscheme "k8s.io/client-go/kubernetes/scheme"
    ctrl "sigs.k8s.io/controller-runtime"
    "sigs.k8s.io/controller-runtime/pkg/healthz"
    "sigs.k8s.io/controller-runtime/pkg/log/zap"

    adapterv1 "github.com/uri-tech/nimble-opti-adapter/api/v1"
    "github.com/uri-tech/nimble-opti-adapter/internal/controller"
    //+kubebuilder:scaffold:imports
)

var (
    scheme = runtime.NewScheme()
    setupLog = ctrl.Log.WithName("setup")
)

func init() {
    // debug
    klog.InfoS("debug - init")

    utilruntime.Must(clientgoscheme.AddToScheme(scheme))

    utilruntime.Must(adapterv1.AddToScheme(scheme))
    //+kubebuilder:scaffold:scheme
}

func main() {
    // debug
    klog.InfoS("debug - main")

    // Define command-line flags.
    var metricsAddr string
    var enableLeaderElection bool
    var probeAddr string
```

```

    flag.StringVar(&metricsAddr, "metrics-bind-address", ":8080", "The
address the metric endpoint binds to.")
    flag.StringVar(&probeAddr, "health-probe-bind-address", ":8081", "The
address the probe endpoint binds to.")
    flag.BoolVar(&enableLeaderElection, "leader-elect", false,
        "Enable leader election for controller manager. "+
        "Enabling this will ensure there is only one active controller
manager.")
    opts := zap.Options{
        Development: true,
    }
    opts.BindFlags(flag.CommandLine)
    flag.Parse()

    // Set the logger for the controller-runtime package.
    ctrl.SetLogger(zap.New(zap.UseFlagOptions(&opts)))

    // Create a new manager to provide shared dependencies and start
components.
    mgr, err := ctrl.NewManager(ctrl.GetConfigOrDie(), ctrl.Options{
        Scheme:             scheme,
        MetricsBindAddress:  metricsAddr,
        Port:                9443,
        HealthProbeBindAddress: probeAddr,
        LeaderElection:      enableLeaderElection,
        LeaderElectionID:    "8f24f142.uri-tech.github.io",
        // LeaderElectionReleaseOnCancel defines if the leader should step
down voluntarily
        // when the Manager ends. This requires the binary to immediately
end when the
        // Manager is stopped, otherwise, this setting is unsafe. Setting
this significantly
        // speeds up voluntary leader transitions as the new leader don't
have to wait
        // LeaseDuration time first.
        //
        // In the default scaffold provided, the program ends immediately
after
        // the manager stops, so would be fine to enable this option.
However,
        // if you are doing or is intended to do any operation such as
perform cleanups
        // after the manager stops then its usage might be unsafe.
        // LeaderElectionReleaseOnCancel: true,
    })
    if err != nil {
        setupLog.Error(err, "unable to start manager")
        os.Exit(1)
    }

    // Initialize the Kubernetes client.
    kubernetesClient := kubernetes.NewForConfigOrDie(mgr.GetConfig())

    // Initialize the IngressWatcher.

```

```

    ingressWatcher := &controller.IngressWatcher{
        Client: kubernetesClient,
    }

    // Pass the KubernetesClient and IngressWatcher to the
    NimbleOptiReconciler.
    if err = (&controller.NimbleOptiReconciler{
        Client: mgr.GetClient(),
        Scheme: mgr.GetScheme(),
        KubernetesClient: kubernetesClient,
        IngressWatcher: ingressWatcher,
    }).SetupWithManager(mgr); err != nil {
        setupLog.Error(err, "unable to create controller", "controller",
"NimbleOpti")
        os.Exit(1)
    }
    //+kubebuilder:scaffold:builder

    // if err = (&controller.NimbleOptiReconciler{
    // Client: mgr.GetClient(),
    // Scheme: mgr.GetScheme(),
    // }).SetupWithManager(mgr); err != nil {
    // setupLog.Error(err, "unable to create controller", "controller",
"NimbleOpti")
    // os.Exit(1)
    // }
    // //+kubebuilder:scaffold:builder

    // Add a health check to the manager.
    if err := mgr.AddHealthzCheck("healthz", healthz.Ping); err != nil {
        setupLog.Error(err, "unable to set up health check")
        os.Exit(1)
    }

    // Add a readiness check to the manager.
    if err := mgr.AddReadyzCheck("readyz", healthz.Ping); err != nil {
        setupLog.Error(err, "unable to set up ready check")
        os.Exit(1)
    }

    // Start the manager and listen for the termination signal.
    setupLog.Info("starting manager")
    if err := mgr.Start(ctrl.SetupSignalHandler()); err != nil {
        setupLog.Error(err, "problem running manager")
        os.Exit(1)
    }
}

```

./api/v1/nimbleopti_types.go:

```
package v1
```

```

import (
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
)

// EDIT THIS FILE!  THIS IS SCAFFOLDING FOR YOU TO OWN!
// NOTE: json tags are required.  Any new fields you add must have json
// tags for the fields to be serialized.

// NimbleOptiSpec defines the desired state of NimbleOpti
type NimbleOptiSpec struct {
    // TargetNamespace is the namespace where the operator should manage
    // certificates
    // +kubebuilder:validation:MinLength=1
    TargetNamespace string `json:"targetNamespace"`

    // CertificateRenewalThreshold is the waiting time (in days) before
    // the certificate expires to trigger renewal
    // +kubebuilder:validation:Minimum=1
    CertificateRenewalThreshold int `json:"certificateRenewalThreshold"`

    // AnnotationRemovalDelay is the delay (in seconds) after removing the
    // "nginx.ingress.kubernetes.io/backend-protocol: HTTPS" annotation before
    // re-adding it
    // +kubebuilder:validation:Minimum=1
    AnnotationRemovalDelay int `json:"annotationRemovalDelay"`

    // RenewalCheckInterval is the interval (in minutes) for checking
    // certificate renewals
    // +kubebuilder:validation:Minimum=1
    RenewalCheckInterval int `json:"renewalCheckInterval"`
}

// NimbleOptiStatus defines the observed state of NimbleOpti
type NimbleOptiStatus struct {
    // Conditions are the conditions for this resource.
    Conditions []metav1.Condition `json:"conditions,omitempty"`

    // IngressPathsForRenewal is a list of ingress paths for which
    // certificates need to be renewed.
    IngressPathsForRenewal []string
    `json:"ingressPathsForRenewal,omitempty"`
}

//+kubebuilder:object:root=true
//+kubebuilder:subresource:status

// NimbleOpti is the Schema for the nimbleoptiadapters API
type NimbleOpti struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec NimbleOptiSpec `json:"spec,omitempty"`
    Status NimbleOptiStatus `json:"status,omitempty"`
}

```

```
//+kubebuilder:object:root=true

// NimbleOptiList contains a list of NimbleOpti
type NimbleOptiList struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ListMeta `json:"metadata,omitempty"`
    Items           []NimbleOpti `json:"items"`
}

func init() {
    SchemeBuilder.Register(&NimbleOpti{}, &NimbleOptiList{})
}
```

For adding control of ingress in the ServiceAccount we modify `./config/rbac/auth_proxy_role.yaml`:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    app.kubernetes.io/name: clusterrole
    app.kubernetes.io/instance: proxy-role
    app.kubernetes.io/component: kube-rbac-proxy
    app.kubernetes.io/created-by: nimble-opti-adapter
    app.kubernetes.io/part-of: nimble-opti-adapter
    app.kubernetes.io/managed-by: kustomize
  name: proxy-role
rules:
- apiGroups:
  - authentication.k8s.io
  resources:
  - tokenreviews
  verbs:
  - create
- apiGroups:
  - authorization.k8s.io
  resources:
  - subjectaccessreviews
  verbs:
  - create
- apiGroups:
  - networking.k8s.io
  resources:
  - ingresses
  verbs:
  - get
  - list
  - watch
  - create
  - delete
  - patch
  - update
```

8. Generate CRD manifests:

When you modify the API definitions, you should regenerate CRD manifests using kubebuilder or operator-sdk tools (whichever you are using)

```
make manifests
make build
```

9. Build and Test the Operator 🚧

If we are using the default Kubebuilder setup:

1. **make install**: This will install the CRDs into the cluster.
2. **make run**: This will run the operator outside of the cluster (for testing purposes).
3. **make docker-build**: This will build the docker image.
4. **make docker-push**: This will push the docker image to the specified registry.
5. **make deploy**: This will deploy the operator into the cluster.

First, let's build our operator's Docker image:

```
docker build -t nimbleopti/nimble-opti-adapter:latest .
docker push nimbleopti/nimble-opti-adapter:latest
```

Crafting a Helm Chart for the Nimble-Opti-Adapter Operator 🚀

1. Pre-requisites 📋

Ensure you've got these tools in your developer's toolbox:

- Helm 3.x installed
- **kubectl** (to interact with your cluster)
- Access to a Kubernetes cluster

2. Initializing the Helm Chart 🛠️

To create a new chart, we'll be using the **helm create** command:

```
cd /mnt/c/project/nimble-opti-adapter
helm create nimble-opti-adapter
```

This will generate a new directory called **nimble-opti-adapter** with the structure of a standard Helm chart.

3. Understanding the Chart Structure 🗺️

Navigate to your `nimble-opti-adapter` directory, and you'll see a structure similar to:

```
nimble-opti-adapter/
├── charts/ # Dependencies (other charts this chart depends on)
├── templates/ # Kubernetes YAML templates with variables
│   ├── deployment.yaml
│   ├── service.yaml
│   ├── _helpers.tpl # Template helpers and definitions
│   └── ...
├── values.yaml # Default configuration values for the templates
└── Chart.yaml # Information about the chart
```

4. Tailoring the Chart for Our Operator 🛠️

1. **Chart.yaml** 📄 : Begin by updating the `Chart.yaml` with relevant details like the version, description, and maintainers.
2. **values.yaml** 📄 : This file holds the default values for your Helm chart. Update it with parameters specific to the `nimble-opti-adapter` operator, like image repository, tag, resources, etc.
3. **Templates** 🗺️ : In the `templates/` directory, customize the existing templates or add new ones. Given this is an operator, you'd likely have:
 - A Deployment for the operator
 - CustomResourceDefinitions (CRDs) the operator will manage
 - RBAC roles and role bindings
 - Any additional resources the operator might need

Remember to make use of Helm's templating functions to ensure flexibility and configurability.

4. complete chart structure :

```
nimble-opti-adapter/ # Root directory of your Helm chart.
├── Chart.yaml # Contains metadata about the Helm chart
like its version, description, etc.
├── README.md # Documentation for how to use the chart.
├── charts # Directory for Helm chart dependencies
(other charts that this chart relies upon).
├── crds # Custom Resource Definitions (CRDs) that
should be installed before any other resources.
│   └── nimbleopti.yaml # The CRD definition for NimbleOpti.
├── templates # Kubernetes resource templates.
│   └── NOTES.txt # Notes that will be displayed to the
```

```
user after helm install or helm upgrade.
|   |— _helpers.tpl           # Define helper methods and variables for
templating.
|   |— deployment.yaml       # Template for the Deployment resource.
Can be used to deploy applications, controllers, etc.
|   |— nimbleopti-instance.yaml # Template for an instance of the
NimbleOpti CRD.
|   |— values.yaml           # Default configuration values that the
templates will use. Can be overridden by the user during install/upgrade.
```

5. Linting & Testing

Before deploying, it's essential to ensure your chart is free from errors:

```
helm lint nimble-opti-adapter
```

It should return:

```
==> Linting nimble-opti-adapter
[INFO] Chart.yaml: icon is recommended
1 chart(s) linted, 0 chart(s) failed
```

To test your Helm chart's rendering:

```
helm install nimble-opti-adapter ./nimble-opti-adapter --dry-run --debug
```

9. Deploying the Operator

Option 1 - using helm:

install the CRDs and deploy the operator:

```
helm install nimble-opti-adapter ./nimble-opti-adapter
```

And monitor the release:

```
helm list
```

or using this guide: [install using helm](#)

Option 2 - directly:

Deploy the CRDs:

```
make install
```

Now, deploy the operator in your cluster:

```
make deploy IMG=nimbleopti/nimble-opti-adapter:latest
```

10. Using github action to auto build and push the image:

create the dir:

```
mkdir -p /mnt/c/project/nimble-opti-adapter/.github/workflows
cd /mnt/c/project/nimble-opti-adapter/.github/workflows
```

create the file `NimbleOpti_CI.yml`:

```
name: NimbleOpti_CI

on:
  workflow_dispatch:
  push:
    branches:
      - "main"
    # pull_request:
    #   branches: [ "main" ]
    # paths:
    #   # - "controllers/**"
    #   # - "Dockerfile"
    #   # - "main.go"
    #   - "*"
    #   # Exclude README.md from triggering the workflow
    paths-ignore:
      - "README.md"
jobs:
  build:
    runs-on: ubuntu-latest

    steps:
      - name: Check the repo.
        uses: actions/checkout@v3

      - name: Login to docker hub registry.
        run: |
```

```

    docker login docker.io -u ${ secrets.DOCKER_USERNAME } -p ${ secrets.DOCKER_PASSWORD }

    # - name: Build and push "NimbleOpti".
    #   run: |
    #     IMAGE_TAG=$(date +%s)
    #     docker build ./NimbleOpti --file ./NimbleOpti/Dockerfile --tag
nimbleopti/nimble-opti-adapter:$IMAGE_TAG --tag nimbleopti/nimble-opti-
adapter:latest
    #     docker push nimbleopti/nimble-opti-adapter:$IMAGE_TAG
    #     docker push nimbleopti/nimble-opti-adapter:latest

- name: Build and push "NimbleOpti".
  run: |
    IMAGE_TAG=$(date +%s)
    docker run --rm --privileged tonistiigi/binfmt --install all
    docker buildx create --use
    DOCKER_TARGET_PLATFORM="linux/arm64,linux/amd64"
    DOCKER_USERNAME="nimbleopti"
    DOCKER_IMAGE_NAME="nimble-opti-adapter"
    docker buildx build . --platform $DOCKER_TARGET_PLATFORM \
    --build-arg BUILD_DATE=$(date -u +'%Y-%m-%dT%H:%M:%SZ') \
    --build-arg VCS_REF=$(git rev-parse --short HEAD) \
    --build-arg VERSION=0.1 \
    --tag $DOCKER_USERNAME/$DOCKER_IMAGE_NAME:$IMAGE_TAG --tag
$DOCKER_USERNAME/$DOCKER_IMAGE_NAME:latest \
    --file ./Dockerfile \
    --output type=image,push=true

```

11. Watch Your Operator in Action 🍿

Create instances of your `NimbleOpti` custom resource and watch your operator ensure the desired state!

12. Celebrate! 🎉

Congratulations! We just finish building a Kubernetes operator using `kubebuilder`! 🥳

docker info:

```

nimbleopti
&mQ6S2L2Yg57$xY

```

problem fix:

```

go clean -modcache
go mod tidy
go mod download
go get sigs.k8s.io/controller-runtime@latest

```

scrpit - create local env and install operator:

```
#!/bin/bash
set -e

DOCKER_USERNAME=${DOCKER_USERNAME:-nimbleopti}
DOCKER_IMAGE_NAME=${DOCKER_IMAGE_NAME:-${DOCKER_USERNAME}/nimble-opti-adapter:latest}
CERT_MANAGER_VERSION=${CERT_MANAGER_VERSION:-v1.11.0}
SLEEP_TIME=${SLEEP_TIME:-1}

# echo "Login to docker..."
# echo $DOCKER_PASSWORD | docker login -u $DOCKER_USERNAME --password-stdin

echo "Deleting Minikube..."
minikube delete

echo "Starting Minikube..."
minikube start

echo "Making manifests..."
make manifests

echo "Installing..."
make install

echo "Building Docker image..."
docker build -t $DOCKER_IMAGE_NAME .

echo "Pushing Docker image..."
docker push $DOCKER_IMAGE_NAME

echo "Deploying..."
make deploy IMG=$DOCKER_IMAGE_NAME

echo "Patching deployment..."
kubectl patch deployment nimble-opti-adapter-controller-manager -n nimble-opti-adapter-system -p '{"spec":{"template":{"spec":{"containers":[{"name":"kube-rbac-proxy","imagePullPolicy":"Always"}, {"name":"manager","imagePullPolicy":"Always"}]}}}}'

echo "Setting Minikube context..."
kubectl config use-context minikube

echo "Adding helm repo..."
helm repo add jetstack https://charts.jetstack.io

echo "Updating helm repo..."
helm repo update
```

```

echo "Installing cert-manager..."
helm install \
  cert-manager jetstack/cert-manager \
  --namespace cert-manager \
  --create-namespace \
  --version $CERT_MANAGER_VERSION \
  --set installCRDs=true \
  --wait

echo "Enabling Minikube ingress..."
minikube addons enable ingress

echo "Starting Minikube dashboard..."
minikube dashboard --url &

# Sleep for 1 second to allow the dashboard to start
sleep $SLEEP_TIME

echo "Setup complete."

```

scrpit - update operator:

```

#!/bin/bash

# causes the shell to exit if any invoked command exits with a non-zero
status
set -e

DOCKER_USERNAME=${DOCKER_USERNAME:-nimbleopti}
DOCKER_IMAGE_NAME=${DOCKER_IMAGE_NAME:-${DOCKER_USERNAME}/nimble-opti-
adapter:latest}

echo "Patching deployment..."
kubectl patch deployment nimble-opti-adapter-controller-manager -n nimble-
opti-adapter-system -p '{"spec":{"template":{"spec":{"containers":
[{"name":"kube-rbac-proxy","imagePullPolicy":"Always"},
{"name":"manager","imagePullPolicy":"Always"}]}}}}'

echo "Making manifests..."
make manifests

echo "Installing..."
make install

echo "Building Docker image..."
docker build -t $DOCKER_IMAGE_NAME .

echo "Pushing Docker image..."
docker push $DOCKER_IMAGE_NAME

echo "Deploying..."
make deploy IMG=$DOCKER_IMAGE_NAME

```

```

echo "Rolling out updates..."
kubectl rollout restart deployment/nimble-opti-adapter-controller-manager
-n nimble-opti-adapter-system
kubectl rollout status deployment/nimble-opti-adapter-controller-manager -
n nimble-opti-adapter-system

echo "Patching deployment..."
kubectl patch deployment nimble-opti-adapter-controller-manager -n nimble-
opti-adapter-system -p '{"spec":{"template":{"spec":{"containers":
[{"name":"kube-rbac-proxy","imagePullPolicy":"Always"},
{"name":"manager","imagePullPolicy":"Always"}]}}}}'

echo "Update complete."

```

script - testing the operator:

```

#!/bin/bash

# causes the shell to exit if any invoked command exits with a non-zero
status
set -e

echo "Applying Ingress resource..."

cat <<EOF | kubectl apply -f -
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
  labels:
    nimble.opti.adapter/enabled: "true"
spec:
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port:
              number: 80
EOF

echo "Ingress resource applied successfully."

echo "Applying NimbleOpti resource..."

```

```
cat <<EOF | kubectl apply -f -
apiVersion: adapter.uri-tech.github.io/v1
kind: NimbleOpti
metadata:
  name: example-nimbleopti
spec:
  # Add any fields defined in the NimbleOpti spec here
  # For instance:
  # someField: someValue
EOF

echo "NimbleOpti resource applied successfully."
```